



# **DevOps Shack**

# **HIGHLY AVAILABLE WEB APPLICATION**



## INTRODUCTION

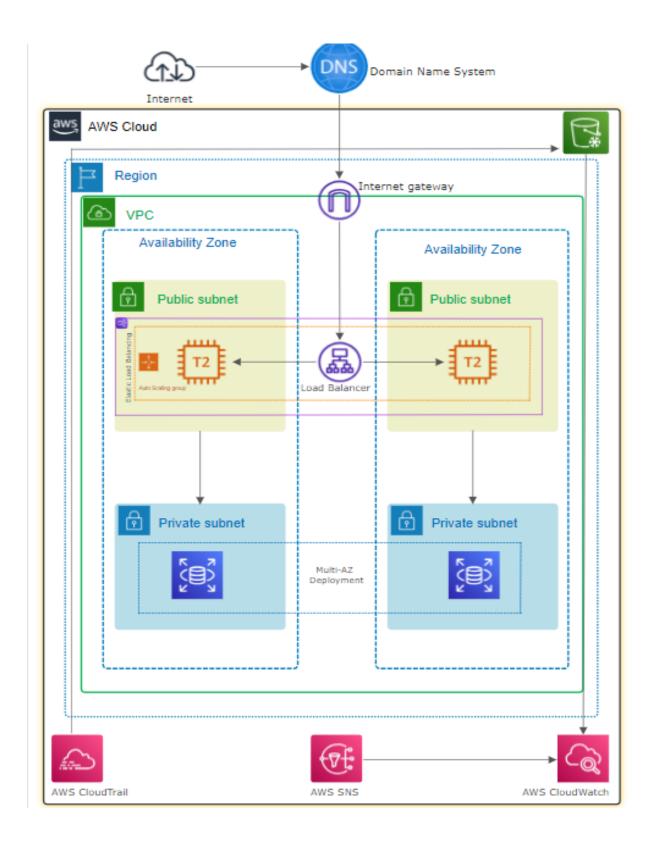
In the dynamic environment of today's web applications, finding uninterrupted service and good performance becomes important. The need for powerful and capable devices has never been more important as businesses strive to meet the ever-increasing demands of their users. This project provides an in-depth look at creating and deploying web applications that can run on Amazon Web Services (AWS).

The importance of versatility cannot be underestimated, especially in the context of web application services, which are the lifeblood of countless businesses. Users expect uninterrupted access to services at all times, so interruptions and interruptions are unacceptable. High availability not only provides a consistent user experience, but also reduces the impact of product failure, network issues, or unexpected disasters.

Using the power of AWS, our goal is to create a design that not only meets the immediate needs of the web, but also adapts and adapts to future initial needs. The main goal of this project is to demonstrate the effectiveness of the web hosting model hosted on AWS. By applying a range of methods that include multi-region availability, load balancing, auto-scaling, and the use of other AWS services, we work to create systems that can be relied upon across projects to differentiate while ensuring service continuity.



# **DATA FLOW DIAGRAMS (DFD)**







## **BACKGROUND**

In the contemporary digital landscape, the relentless availability of web applications is not merely a convenience but a fundamental expectation. Users demand seamless access to services 24/7, placing an emphasis on high availability to ensure a consistently positive user experience. Simultaneously, businesses recognize that downtime can translate into significant financial losses, disrupt operations, and tarnish their reputation.

Amazon Web Services (AWS) emerges as a pivotal solution in meeting the demands of high availability for web applications. With its architecture built on multiple geographically distributed data centers called Availability Zones, AWS provides a foundation for redundancy and fault tolerance. Elastic Load Balancing optimizes the distribution of incoming traffic across instances, preventing the risk of a single point of failure. Auto Scaling, another key feature, enables dynamic adjustments of computing resources based on demand, ensuring optimal performance during peak periods and cost efficiency during troughs.

Managed database services, such as Amazon RDS, contribute to high availability by offering features like automated backups, failover support, and seamless scalability. AWS CloudFront, a content delivery network, enhances user experience by caching and delivering content from edge locations globally, reducing latency and improving responsiveness.

Crucially, AWS's management of a wide array of services, from infrastructure to security, allows businesses to focus on application development rather than the complexities of infrastructure management. The robust security features, coupled with compliance certifications, further reinforce the reliability and resilience of applications hosted on AWS.



# **REQUIREMENTS**

## **System Architecture**

#### 1. Amazon EC2 Instances:

- Deploy two EC2 instances in a public subnet inside a VPC
- •
- Set up Auto Scaling Groups to dynamically scale instances based on demand.
- Ensure instances are distributed across different Availability Zones for enhanced fault tolerance.

### 2. Load Balancer:

- Implement a load balancer to distribute user traffic evenly across EC2 instances.
- Enhance system fault tolerance and availability through effective load balancing.

### 3. Amazon RDS:

- Establish two RDS instances in a private subnet.
- Choose different Availability Zones for each RDS instance to ensure redundancy.
- Set up database replication for real-time data synchronization.
- Enable multi-AZ deployment for improved fault tolerance and reliability.

## **Security Considerations**

### 1. Networking:

 Utilize private and public subnets to segregate resources for improved security.





 Implement security groups and Network ACLs to control inbound and outbound traffic effectively.

## 2. Encryption:

Enable encryption for data in transit and at rest within RDS instances.

## **Monitoring and Logging**

### 1. CloudWatch:

- Integrate CloudWatch to monitor system metrics, logs, and performance.
- Set up alarms to proactively respond to issues and maintain optimal performance.

### 2. CloudTrail:

- Implement AWS CloudTrail to capture API calls for auditing and compliance.
- Store logs securely in an S3 bucket for centralized access.

#### **Notification Services**

### 1. Amazon SNS:

- Integrate Amazon SNS for sending notifications.
- Configure SNS to send timely alerts based on CloudWatch alarms, ensuring prompt responses to system issues.

## **Scalability and Redundancy**

### 1. Auto Scaling:

- Implement Auto Scaling Groups to allow automatic scaling of EC2 instances based on varying demand.
- 2. Redundancy:
- Ensure multi-AZ deployment for both EC2 instances and RDS, contributing to system redundancy and high availability.



### **OVERVIEW**

In this project, the goal is to construct a robust and highly available web application leveraging Amazon Web Services (AWS). The architecture is meticulously designed to ensure uninterrupted service, scalability, and reliability, with a focus on the following key components:

#### 1. Amazon EC2 Instances:

Two EC2 instances are deployed across different Availability Zones in a public subnet.

Auto Scaling Groups are implemented to dynamically scale resources based on demand, ensuring optimal performance under varying workloads.

### 2. Load Balancer:

A load balancer is employed to evenly distribute incoming traffic across the EC2 instances, enhancing fault tolerance and availability.

#### 3. Amazon RDS:

Two RDS instances reside in a private subnet, each in a different Availability Zone.

Database replication is implemented for real-time synchronization, ensuring data availability and reliability.

Multi-AZ deployment enhances fault tolerance, providing continuous service even in the face of potential disruptions.

## 4. Security Considerations:

Networking is optimized through private and public subnets, segregating resources for improved security.





Encryption mechanisms are implemented to secure data both in transit and at rest within RDS instances.

# 5. Monitoring and Logging:

CloudWatch is integrated for real-time monitoring of system metrics, logs, and overall performance.

AWS CloudTrail captures API calls, storing logs securely in an S3 bucket for auditing and compliance.

### 6. Notification Services:

Amazon Simple Notification Service (SNS) is utilized to send notifications, configured to alert based on CloudWatch alarms, ensuring swift responses to system issues.

Scalability and Redundancy:

Auto Scaling Groups and multi-AZ deployment contribute to system scalability and redundancy, accommodating varying workloads and enhancing fault tolerance





### **DESCRIPTION**

## **1.Amazon EC2 Instances:**

**Role**: Serve as the compute infrastructure hosting the web application.

### **Functionality:**

Two EC2 instances are deployed, each in a different Availability Zone for fault tolerance.

Auto Scaling Groups are employed for dynamic scaling based on varying workloads.

Instances run the Ubuntu operating system and host the web application.

### 2. Load Balancer:

**Role:** Distribute incoming traffic across multiple EC2 instances.

### **Functionality:**

Ensures even distribution of traffic, preventing overload on individual instances.

Enhances fault tolerance by redirecting traffic away from unhealthy instances.

Facilitates seamless scaling by adjusting to changes in the number of instances.

## 3. Amazon RDS:

Role: Host and manage the relational database for the web application.

### **Functionality:**

Two RDS instances are deployed in different Availability Zones to ensure redundancy.

Database replication is implemented for real-time synchronization between primary and standby instances.

Multi-AZ deployment enhances fault tolerance by automatically failing over to the standby instance in case of a primary instance failure.

Supports MySQL or another relational database engine.

### 4. Security Considerations:

**Role:** Ensure the security and integrity of the entire system.



## **Functionality:**

Utilizes private and public subnets for resource segregation.

Implements security groups to control inbound and outbound traffic to EC2 instances.

Network ACLs are configured to add an additional layer of security to the subnet.

Encryption is applied for data in transit and at rest within RDS instances, ensuring data security.

## 5. Monitoring and Logging:

Role: Provide real-time insights into system performance and activity.

## **Functionality:**

CloudWatch is integrated for monitoring EC2 instances, load balancer, and RDS metrics.

Alarms are configured to notify administrators of predefined threshold breaches.

AWS CloudTrail captures API calls, providing a detailed history of changes and activities within the AWS environment.

Logs are stored securely in an S3 bucket for auditing and compliance.

## 6. Notification Services (Amazon SNS):

**Role:** Notify administrators of critical events or issues.

### **Functionality:**

Amazon SNS is configured to send notifications based on CloudWatch alarms.

Alerts are sent to predefined recipients, ensuring timely responses to system issues.

## 7. Scalability and Redundancy:

**Role:** Ensure the system can handle varying workloads and maintain continuous service availability.

## **Functionality:**





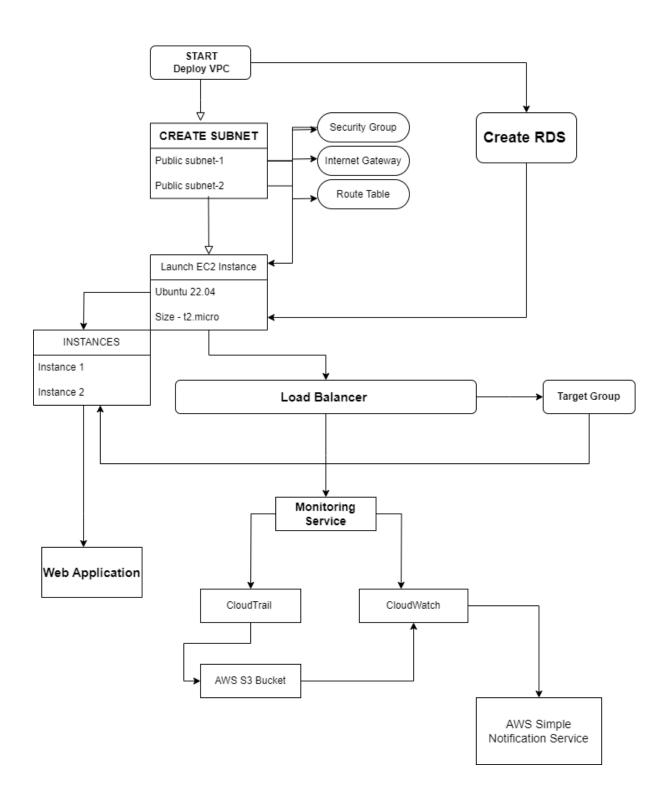
Auto Scaling Groups dynamically adjust the number of EC2 instances based on demand.

Multi-AZ deployment of EC2 instances and RDS provides redundancy and fault tolerance, minimizing the impact of potential failures.





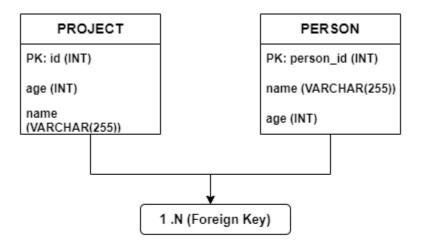
# FLOWCHART (MODULES & SUB-MODULES)







# **DATA DESCRIPTION (E-R DIAGRAMS)**



PROJECT Table Entries		
ID	age	Name
1	25	Shubham
2	30	Raj
3	22	Abhishek
4	28	Riya
5	35	Helen
6	35	Rayan
7	40	Sayan
8	32	Aman
9	24	Reha
10	32	Mehta





## **SAMPLE CODE**

# Script 1: RDS connection from ec2

sudo su

sudo apt update

sudo apt install mysql-server

sudo systemctl start mysql

sudo systemctl enable mysql

mysql -h mysql.crlw2n0fdtpd.us-east-1.rds.amazonaws.com -u admin -p

Enter password - admin123



## Script 2: RDS dummy data insertion

```
CREATE TABLE PROJECT (
  id INT PRIMARY KEY,
  age INT,
  name VARCHAR(255)
);
INSERT INTO PROJECT (id, age, name) VALUES
(1, 25, 'Shubham'),
(2, 30, 'Raj'),
(3, 22, 'Abhishek'),
(4, 28, 'Riya'),
(5, 35, 'Helen');
SELECT * FROM PROJECT;
INSERT INTO PROJECT (id, age, name) VALUES
(6, 35, 'Rayan'),
(7, 40, 'Sayan'),
(8, 32, 'Aman'),
(9, 24, 'Reha'),
(10, 32, 'Mehta');
```





## Script 3: Apache server fetches data from RDS

```
#!/bin/bash
# Database configuration
host="mysql.crlw2n0fdtpd.us-east-1.rds.amazonaws.com"
username="admin"
password="admin123"
database="mydatabase"
table="PROJECT"
# Install required packages
sudo apt update
sudo apt install -y apache2 php libapache2-mod-php php-mysql
# Create a PHP script to fetch data from the database
sudo tee /var/www/html/index.php >/dev/null <<EOF
<?php
\$conn = new mysqli("$host", "$username", "$password", "$database");
if (\$conn->connect_error) {
  die("Connection failed: ".\$conn->connect error);
\$sql = "SELECT * FROM $table";
\ result = \ conn->query(\sql);
if (\$result->num rows > 0) {
  echo "IDName";
```





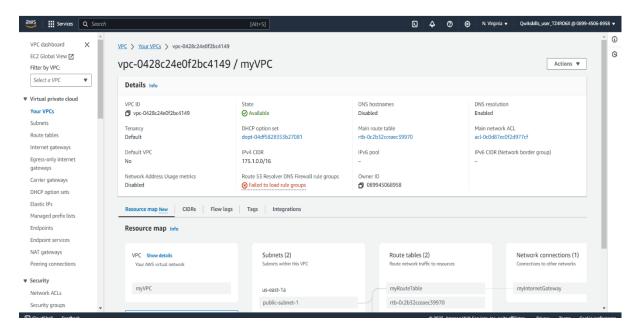
```
while(\$row = \$result->fetch assoc()) {
    echo "" . \$row["id"] . "" . \$row["name"] . "";
  echo "";
} else {
  echo "0 results";
\$conn->close();
?>
EOF
# Configure Apache to serve PHP files
sudo sed -i "s/index.html/index.php/g" /etc/apache2/mods-enabled/dir.conf
# Restart Apache
sudo systemctl restart apache2
# Display public IP
echo "Web application is now accessible at: http://$(curl -s
http://checkip.amazonaws.com)/"
```



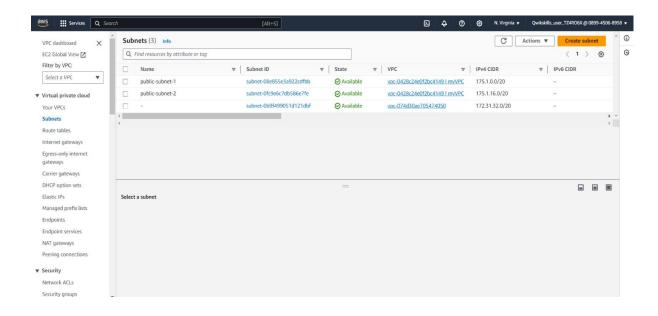


### **SCREENSHOTS**

## Vpc



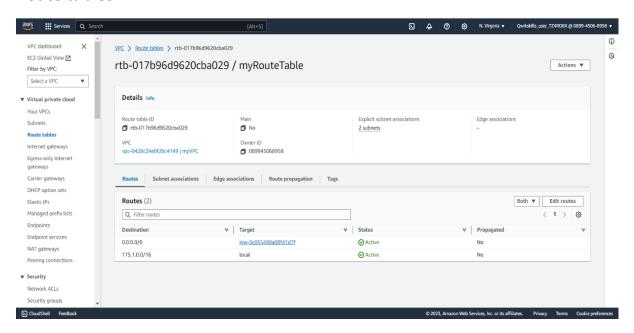
### **Subnets**



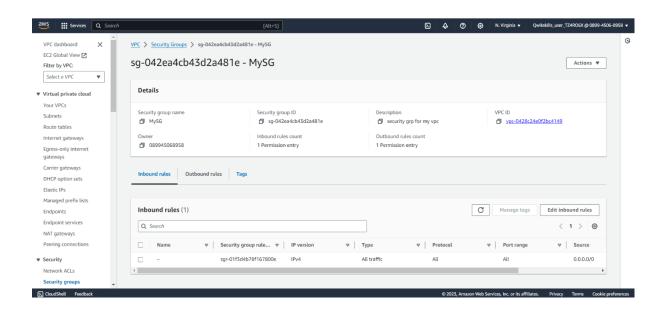




### **Route tables**



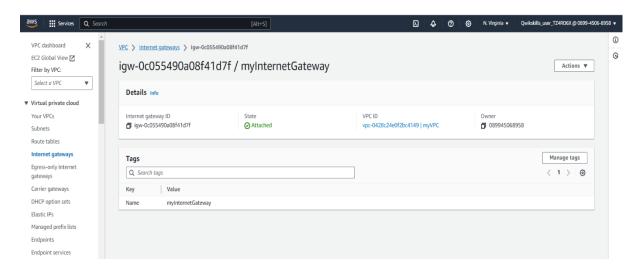
## **Security group**



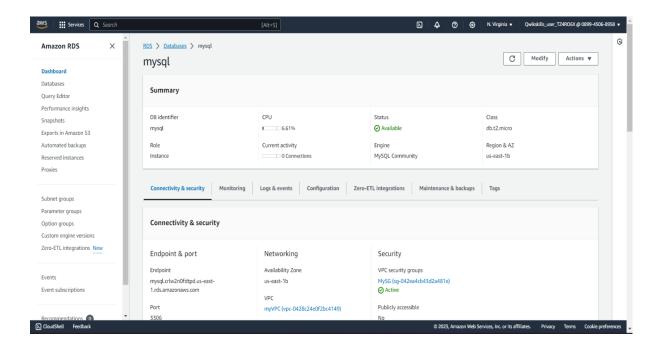




### Internet gateway



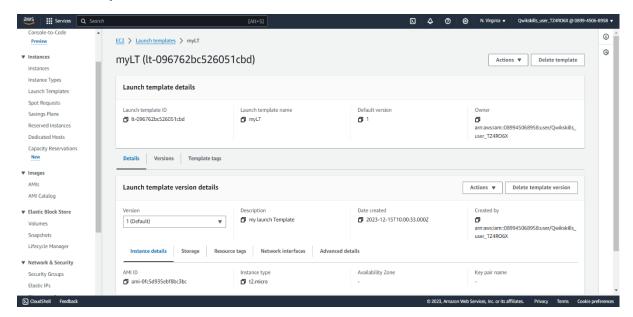
### **RDS** instances



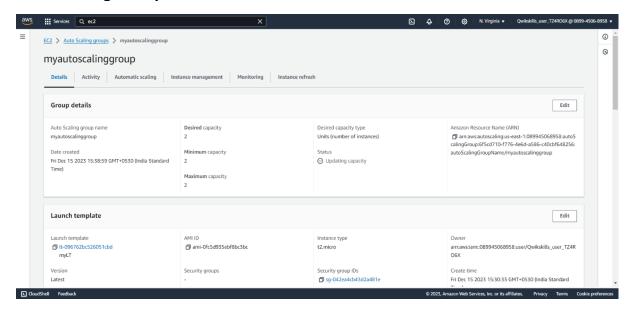




## **Launch Template**



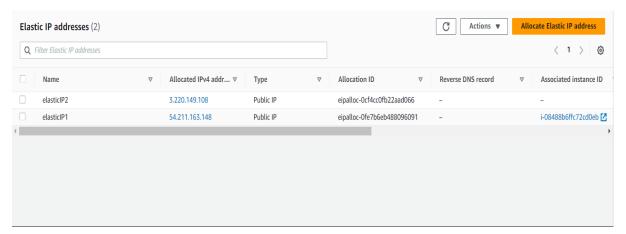
# **Auto Scaling Group**



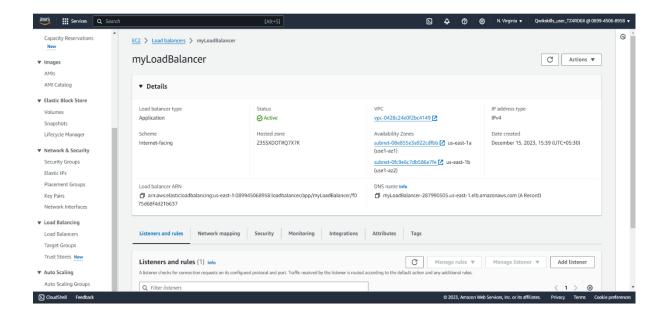




## **Elastic Ips**



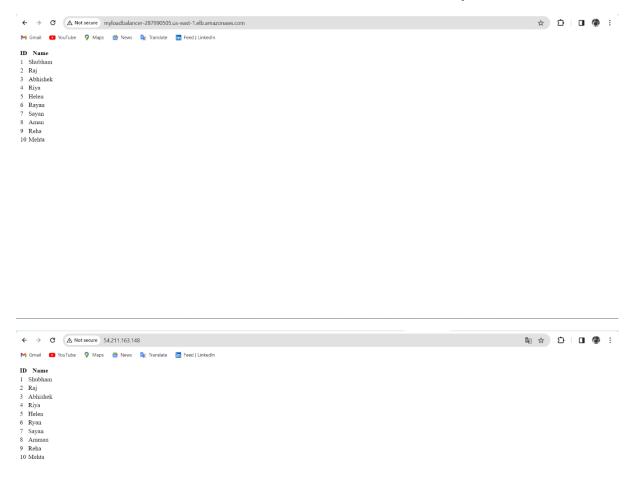
### **Load Balancer**







# DNS result, and result from both EC2 Instances with their public IP

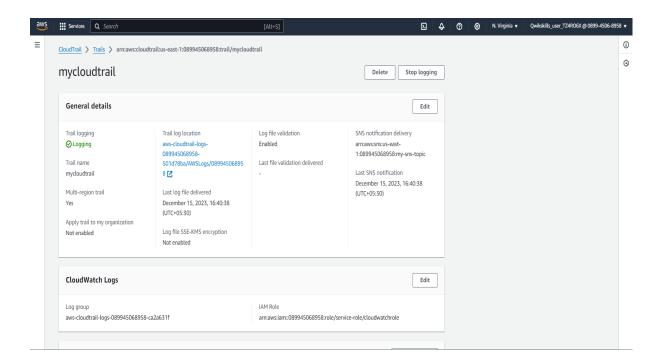








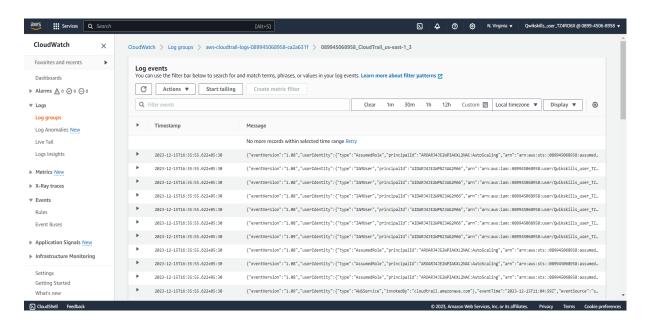
## cloudTrail



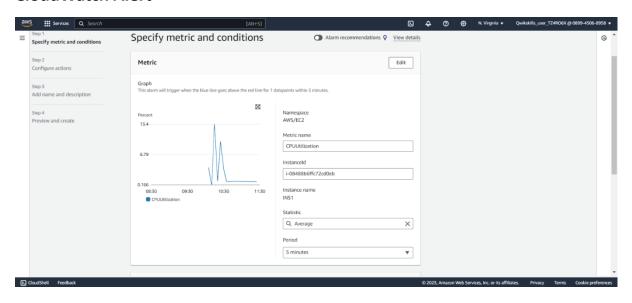




### CloudWatch



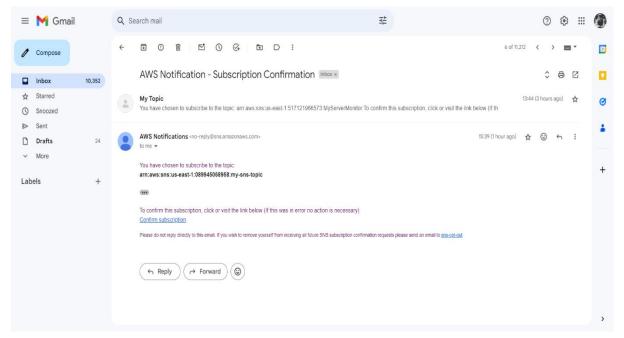
### **CloudWatch Alert**



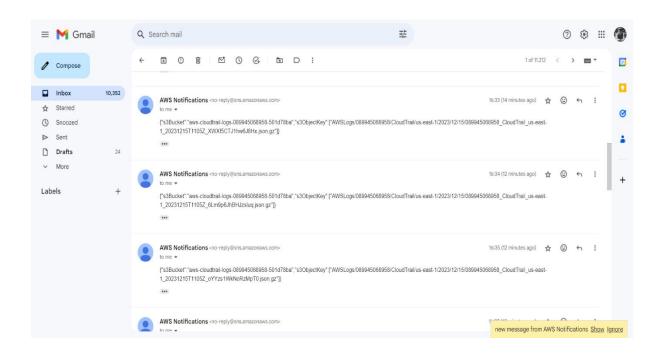




## **Email for SNS subscription**



### **SNS Notification on Email**







### CONCLUSION

In culmination, the design and implementation of the highly available web application on AWS represent a significant achievement in building a robust and scalable infrastructure. The project's success is underscored by the seamless accessibility achieved through both the load balancer DNS name and public IPs, showcasing the efficacy of load balancing and redundancy mechanisms.

Key components of the project's architecture, including the creation of a VPC (175.1.0.0/16) with two public subnets (175.1.0.0/20 and 175.1.16.0/20) strategically placed in two availability zones (us-east1a and us-east1b), Amazon RDS for efficient database management, and strategically deployed EC2 instances across different availability zones, collectively contribute to a resilient framework capable of handling diverse workloads while ensuring uninterrupted accessibility.

The incorporation of an auto-scaling group and load balancer further enhances the project's scalability and fault tolerance, ensuring a seamless user experience even under varying traffic conditions.

The integration of monitoring services, including CloudWatch, CloudTrail, and SNS, enhances the project's overall resilience. Real-time visibility into metrics and logs, coupled with proactive notification mechanisms, adds a layer of observability critical for maintaining system health and addressing potential issues promptly.

As the project journey concludes, it is valuable to reflect on the lessons learned, challenges overcome, and areas poised for future enhancement

In summary, the project not only fulfills the outlined requirements for a highly available web application on AWS but also exemplifies a commitment to best practices in cloud architecture. Congratulations on reaching this milestone, and may the lessons learned pave the way for continued success in future endeavors.



### **SALIENT FEATURES**

### 1. Highly Available Infrastructure:

The project establishes a VPC (175.1.0.0/16) with two public subnets (175.1.0.0/20 and 175.1.16.0/20) distributed across distinct availability zones (us-east1a and us-east1b). This design ensures redundancy and fault tolerance, contributing to the high availability of the web application.

## 2. Elastic Compute Capacity:

Leveraging Amazon EC2 instances and an auto-scaling group, the project dynamically adjusts compute capacity based on demand. This elasticity ensures optimal performance even during peak traffic, enhancing the scalability of the application.

## 3. Managed Database Service:

Amazon RDS is employed for efficient database management. This managed service automates routine administrative tasks, such as backups and software patching, allowing focus on application development rather than database maintenance.

### 4. Load Balancing for Even Workload Distribution:

The implementation of a load balancer intelligently distributes incoming traffic across multiple EC2 instances, preventing overloading of individual instances. This load balancing mechanism enhances the overall performance and reliability of the web application.

## 5. Automated Script for Quick Deployment:

The project includes a Bash script that automates the deployment of the web application, streamlining the setup process for reproducibility and ease of deployment in various environments.





### 6. Monitoring and Alerting System:

Integration with AWS monitoring services, including CloudWatch, CloudTrail, and SNS, provides real-time insights into system health, performance metrics, and logs. Proactive alerting ensures timely responses to potential issues, contributing to the robustness of the application.

## 7. Public and Load Balancer Accessibility:

The web application is made accessible through both public IP addresses and a load balancer DNS name, showcasing effective load balancing and redundancy measures. This dual accessibility enhances the availability and reliability of the application.

### 8. Security Best Practices:

Adherence to AWS security best practices, including the use of security groups and the implementation of secure communication protocols, ensures the integrity and confidentiality of data transmitted within the application.

### 9. Future-Ready Architecture:

The project's architecture is designed with scalability and adaptability in mind, providing a foundation for future enhancements. The inclusion of a reflective documentation process enables continuous improvement based on lessons learned.

In summary, these salient features collectively define a highly available, scalable, and secure web application on AWS, showcasing a commitment to best practices in cloud architecture



### **FUTURE SCOPE**

As the project progresses, key future development areas emerge:

- Feature Expansion:
- Prioritizing the integration of advanced features like user authentication, real-time collaboration, and sophisticated data analytics to enhance the web application's utility and responsiveness.
  - Global Deployment:
- Strategically extending the project's reach by deploying instances in additional AWS regions to reduce latency, fortify redundancy, and ensure continuous service availability globally.
  - CI/CD Implementation:
- Implementing Continuous Integration and Deployment (CI/CD) pipelines to automate testing, building, and deployment processes. This accelerates release cycles, improves code quality, and seamlessly integrates updates.
  - Microservices Exploration:
- Exploring the adoption of a microservices architecture for scalable and maintainable development. Breaking down the monolithic structure into independent services allows for flexible deployment and agile development practices.
  - Advanced Security Measures:
- Advancing the project's security posture through the integration of AWS WAF, robust encryption practices, and regular security audits. These measures ensure ongoing integrity and confidentiality of user data.

In summary, the future project scope revolves around enriching features, extending global accessibility, optimizing development processes, exploring architectural improvements, and fortifying security measures. This collective effort aims to establish a resilient, feature-rich, and secure web application.