▶ DevOps Shack

StatefulSets in Kubernetes

StatefulSets are a Kubernetes resource for managing stateful applications. Unlike Deployments, StatefulSets maintain the identity of each pod, ensuring predictable, ordered pod creation, deletion, and scaling. They are designed to handle workloads that require persistent storage and stable network identifiers, like databases (e.g., MySQL, MongoDB, etc.).

Key Characteristics of StatefulSets:

- **Stable, Unique Network IDs**: Each pod gets a consistent network identifier, which doesn't change after rescheduling.

- **Ordered, Graceful Deployment and Scaling**: Pods are created in sequence (e.g., pod-0, pod-1, pod-2), and any scaling operation follows the same order.

- **Persistent Storage**: Each pod gets its own PersistentVolume (PV) that is not shared with other pods, ensuring that data persists across pod rescheduling or restarts.

Components of a StatefulSet:

1. **StatefulSet Controller**: Manages the orchestration of the pods in a specific order and ensures that they are deployed in a sequence.

2. **Pods**: The containers that hold the application's processes.

3. **Persistent Volume Claims (PVCs)**: StatefulSet automatically creates one PVC per pod to retain the state of the application.

4. **Headless Service**: Used for stable network identities.

StatefulSet Configuration

Below is an example of a basic StatefulSet YAML configuration for a MySQL application:

```yaml
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: mysql
spec:
  serviceName: "mysql"
  replicas: 3
  selector:
    matchLabels:
      app: mysql
  template:
    metadata:
      labels:
        app: mysql
    spec:
      containers:
      - name: mysql
        image: mysql:5.7
        ports:
        - containerPort: 3306
          name: mysql
        volumeMounts:
        - name: mysql-persistent-storage
          mountPath: /var/lib/mysql
  volumeClaimTemplates:
  - metadata:
```

```
   name: mysql-persistent-storage
 spec:
  accessModes: [ "ReadWriteOnce" ]
  resources:
   requests:
    storage: 1Gi
```

Explanation of Key Parts:

1. **serviceName**: Defines a headless service that ensures each pod in the StatefulSet gets a stable network ID.

2. **replicas**: Specifies the number of pod replicas (in this case, 3 MySQL pods).

3. **volumeClaimTemplates**: Automatically creates PersistentVolumeClaims for each pod.

How StatefulSet Works

StatefulSet ensures that:

- Pods are created one at a time (e.g., pod-0, pod-1, pod-2).

- When scaling up, new pods are added in order (e.g., pod-3, pod-4).

- When scaling down, the last pod is deleted first (e.g., pod-4, then pod-3).

- Each pod gets its own storage via PVCs, ensuring that data is not shared among pods.

StatefulSet vs. Deployment

Feature	StatefulSet	Deployment
Pod Identity	Stable (pod-0, pod-1, pod-2)	Dynamic (random names)
Network Identity	Stable, DNS name (<pod-name>.<service-name>.<namespace>)	Dynamic, based on IP address
Persistent Storage	Unique PVC for each pod	Shared storage or ephemeral
Pod Scaling Order	Ordered (creates pods sequentially)	Unordered (creates pods concurrently)
Pod Termination/Restart	Ordered (terminates pods in reverse order)	Unordered (terminates pods randomly)

Working Example with Persistent Volumes

Consider the following StatefulSet, which provisions a three-node Redis cluster. Each Redis node has its own unique identity and persistent storage.

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: redis
spec:
  serviceName: "redis"
  replicas: 3
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
      - name: redis
        image: redis:6.0
        ports:
        - containerPort: 6379
          name: redis
        volumeMounts:
        - name: redis-storage
          mountPath: /data
  volumeClaimTemplates:
  - metadata:
      name: redis-storage
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 2Gi
```

Benefits of StatefulSets

1. **Stable Network Identity**: Each pod gets a predictable, stable network identity (e.g., redis-0, redis-1). This is crucial for stateful applications like databases that need to connect to a specific instance.

2. **Persistent Storage**: Each pod has its own storage, which ensures that data persists even after a pod is rescheduled or restarted. The PersistentVolumeClaim (PVC) associated with each pod ensures that data is not lost.

3. **Ordered, Graceful Scaling and Updates**: StatefulSet ensures that pods are created, updated, and scaled in a specific order, preventing race conditions or inconsistencies in data.

4. **Seamless Scaling**: You can easily scale stateful applications while ensuring data consistency and stability.

Common Use Cases for StatefulSets

1. **Databases**: Databases like MySQL, PostgreSQL, and Cassandra need stable storage and network identities for replication and sharding.

2. **Distributed Systems**: Distributed applications like Kafka, Redis, and ZooKeeper require stateful storage and stable identities for correct cluster functioning.

3. **Message Brokers**: Stateful message brokers like RabbitMQ and NATS that maintain message state between restarts benefit from StatefulSets.

Scaling StatefulSets

To scale the number of replicas in a StatefulSet:

kubectl scale statefulset redis --replicas=5

Kubernetes will create new pods (redis-3, redis-4) in a sequential manner.

To scale down:

kubectl scale statefulset redis --replicas=2

This will terminate pods in reverse order (redis-4 first, then redis-3).

StatefulSet Limitations

1. **Not Ideal for Stateless Applications**: For purely stateless applications, a Deployment or ReplicaSet is more efficient.

2. **Pod Deletion**: Deleting StatefulSet pods is not straightforward. Deleting a pod with an attached PersistentVolume doesn't free the volume; it must be handled manually.

3. **Storage Management**: StatefulSets rely heavily on the underlying storage infrastructure. If the storage layer is not well configured (e.g., PersistentVolumes), this could lead to data loss.

Best Practices for StatefulSets

1. **Use Persistent Storage**: Always ensure that StatefulSet pods use PersistentVolumeClaims for data persistence.

2. **Headless Services**: Use headless services to ensure stable DNS entries for each pod. This guarantees that the pod's identity is preserved even during rescheduling.

3. **Monitor Scaling**: Since StatefulSets scale one pod at a time, monitor your scaling operations to ensure each pod is healthy before adding/removing another pod.

4. **Test Failure Recovery**: Test the resiliency of StatefulSets by simulating pod/node failures and ensuring your application recovers gracefully.

Conclusion

StatefulSets are essential for managing stateful applications in Kubernetes. They ensure stable network identities, persistent storage, and ordered scaling, making them suitable for databases and distributed systems. By using StatefulSets, applications can recover gracefully from failures and maintain their state across different cluster nodes.